

• Data Preprocessing.

- o Import the Libraries.
- o Importing the dataset.
- o Checking for Null Values.
- o Data Visualization.
- o Outlier Detection
- o Splitting Dependent and Independent variables
- o- Encoding
- o Feature Scaling.
- o Splitting Data into Train and Test.

1.Import the Libraries.

```
In [4]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

2.Importing the dataset.

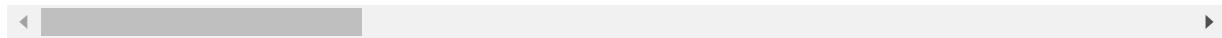
```
In [5]: df=pd.read_csv("Employee-Attrition.csv")
```

In [6]: df.head()

Out[6]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | Education |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|-----------|
| 0 | 41 | Yes | Travel_Rarely | 1102 | Sales | | 1 | 2 |
| 1 | 49 | No | Travel_Frequently | 279 | Research & Development | | 8 | 1 |
| 2 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | | 2 | 2 |
| 3 | 33 | No | Travel_Frequently | 1392 | Research & Development | | 3 | 4 |
| 4 | 27 | No | Travel_Rarely | 591 | Research & Development | | 2 | 1 |

5 rows × 35 columns

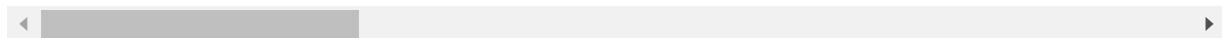


In [7]: df.tail()

Out[7]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | Education |
|------|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|-----------|
| 1465 | 36 | No | Travel_Frequently | 884 | Research & Development | | 23 | 2 |
| 1466 | 39 | No | Travel_Rarely | 613 | Research & Development | | 6 | 1 |
| 1467 | 27 | No | Travel_Rarely | 155 | Research & Development | | 4 | 3 |
| 1468 | 49 | No | Travel_Frequently | 1023 | Sales | | 2 | 3 |
| 1469 | 34 | No | Travel_Rarely | 628 | Research & Development | | 8 | 3 |

5 rows × 35 columns



In [8]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object  
 2   BusinessTravel   1470 non-null    object  
 3   DailyRate        1470 non-null    int64  
 4   Department       1470 non-null    object  
 5   DistanceFromHome 1470 non-null    int64  
 6   Education        1470 non-null    int64  
 7   EducationField   1470 non-null    object  
 8   EmployeeCount    1470 non-null    int64  
 9   EmployeeNumber   1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object  
 12  HourlyRate       1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object  
 16  JobSatisfaction  1470 non-null    int64  
 17  MaritalStatus    1470 non-null    object  
 18  MonthlyIncome    1470 non-null    int64  
 19  MonthlyRate      1470 non-null    int64  
 20  NumCompaniesWorked 1470 non-null    int64  
 21  Over18            1470 non-null    object  
 22  OverTime          1470 non-null    object  
 23  PercentSalaryHike 1470 non-null    int64  
 24  PerformanceRating 1470 non-null    int64  
 25  RelationshipSatisfaction 1470 non-null    int64  
 26  StandardHours    1470 non-null    int64  
 27  StockOptionLevel  1470 non-null    int64  
 28  TotalWorkingYears 1470 non-null    int64  
 29  TrainingTimesLastYear 1470 non-null    int64  
 30  WorkLifeBalance   1470 non-null    int64  
 31  YearsAtCompany   1470 non-null    int64  
 32  YearsInCurrentRole 1470 non-null    int64  
 33  YearsSinceLastPromotion 1470 non-null    int64  
 34  YearsWithCurrManager 1470 non-null    int64  
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

In [9]: df.describe()

Out[9]:

| | Age | DailyRate | DistanceFromHome | Education | EmployeeCount | EmployeeNur |
|-------|-------------|-------------|------------------|-------------|---------------|-------------|
| count | 1470.000000 | 1470.000000 | 1470.000000 | 1470.000000 | 1470.0 | 1470.000000 |
| mean | 36.923810 | 802.485714 | 9.192517 | 2.912925 | 1.0 | 1024.861404 |
| std | 9.135373 | 403.509100 | 8.106864 | 1.024165 | 0.0 | 602.024390 |
| min | 18.000000 | 102.000000 | 1.000000 | 1.000000 | 1.0 | 1.000000 |
| 25% | 30.000000 | 465.000000 | 2.000000 | 2.000000 | 1.0 | 491.250000 |
| 50% | 36.000000 | 802.000000 | 7.000000 | 3.000000 | 1.0 | 1020.500000 |
| 75% | 43.000000 | 1157.000000 | 14.000000 | 4.000000 | 1.0 | 1555.750000 |
| max | 60.000000 | 1499.000000 | 29.000000 | 5.000000 | 1.0 | 2068.000000 |

8 rows × 26 columns



In [10]: df.shape

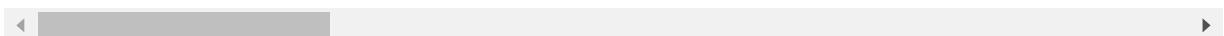
Out[10]: (1470, 35)

In [11]: df.corr()

Out[11]:

| | Age | DailyRate | DistanceFromHome | Education | EmployeeCount | E |
|--------------------------|-----------|-----------|------------------|-----------|---------------|---|
| Age | 1.000000 | 0.010661 | -0.001686 | 0.208034 | NaN | |
| DailyRate | 0.010661 | 1.000000 | -0.004985 | -0.016806 | NaN | |
| DistanceFromHome | -0.001686 | -0.004985 | 1.000000 | 0.021042 | NaN | |
| Education | 0.208034 | -0.016806 | 0.021042 | 1.000000 | NaN | |
| EmployeeCount | NaN | NaN | NaN | NaN | NaN | |
| EmployeeNumber | -0.010145 | -0.050990 | 0.032916 | 0.042070 | NaN | |
| EnvironmentSatisfaction | 0.010146 | 0.018355 | -0.016075 | -0.027128 | NaN | |
| HourlyRate | 0.024287 | 0.023381 | 0.031131 | 0.016775 | NaN | |
| JobInvolvement | 0.029820 | 0.046135 | 0.008783 | 0.042438 | NaN | |
| JobLevel | 0.509604 | 0.002966 | 0.005303 | 0.101589 | NaN | |
| JobSatisfaction | -0.004892 | 0.030571 | -0.003669 | -0.011296 | NaN | |
| MonthlyIncome | 0.497855 | 0.007707 | -0.017014 | 0.094961 | NaN | |
| MonthlyRate | 0.028051 | -0.032182 | 0.027473 | -0.026084 | NaN | |
| NumCompaniesWorked | 0.299635 | 0.038153 | -0.029251 | 0.126317 | NaN | |
| PercentSalaryHike | 0.003634 | 0.022704 | 0.040235 | -0.011111 | NaN | |
| PerformanceRating | 0.001904 | 0.000473 | 0.027110 | -0.024539 | NaN | |
| RelationshipSatisfaction | 0.053535 | 0.007846 | 0.006557 | -0.009118 | NaN | |
| StandardHours | NaN | NaN | NaN | NaN | NaN | |
| StockOptionLevel | 0.037510 | 0.042143 | 0.044872 | 0.018422 | NaN | |
| TotalWorkingYears | 0.680381 | 0.014515 | 0.004628 | 0.148280 | NaN | |
| TrainingTimesLastYear | -0.019621 | 0.002453 | -0.036942 | -0.025100 | NaN | |
| WorkLifeBalance | -0.021490 | -0.037848 | -0.026556 | 0.009819 | NaN | |
| YearsAtCompany | 0.311309 | -0.034055 | 0.009508 | 0.069114 | NaN | |
| YearsInCurrentRole | 0.212901 | 0.009932 | 0.018845 | 0.060236 | NaN | |
| YearsSinceLastPromotion | 0.216513 | -0.033229 | 0.010029 | 0.054254 | NaN | |
| YearsWithCurrManager | 0.202089 | -0.026363 | 0.014406 | 0.069065 | NaN | |

26 rows × 26 columns



3.Checking for Null Values.

In [12]: `df.isnull().any()`

Out[12]:

| | |
|--------------------------|-------|
| Age | False |
| Attrition | False |
| BusinessTravel | False |
| DailyRate | False |
| Department | False |
| DistanceFromHome | False |
| Education | False |
| EducationField | False |
| EmployeeCount | False |
| EmployeeNumber | False |
| EnvironmentSatisfaction | False |
| Gender | False |
| HourlyRate | False |
| JobInvolvement | False |
| JobLevel | False |
| JobRole | False |
| JobSatisfaction | False |
| MaritalStatus | False |
| MonthlyIncome | False |
| MonthlyRate | False |
| NumCompaniesWorked | False |
| Over18 | False |
| Overtime | False |
| PercentSalaryHike | False |
| PerformanceRating | False |
| RelationshipSatisfaction | False |
| StandardHours | False |
| StockOptionLevel | False |
| TotalWorkingYears | False |
| TrainingTimesLastYear | False |
| WorkLifeBalance | False |
| YearsAtCompany | False |
| YearsInCurrentRole | False |
| YearsSinceLastPromotion | False |
| YearsWithCurrManager | False |
| dtype: | bool |

```
In [13]: df.isnull().sum()
```

```
Out[13]: Age          0  
Attrition      0  
BusinessTravel  0  
DailyRate       0  
Department      0  
DistanceFromHome 0  
Education        0  
EducationField    0  
EmployeeCount     0  
EmployeeNumber    0  
EnvironmentSatisfaction 0  
Gender          0  
HourlyRate       0  
JobInvolvement    0  
JobLevel         0  
JobRole          0  
JobSatisfaction   0  
MaritalStatus     0  
MonthlyIncome      0  
MonthlyRate        0  
NumCompaniesWorked 0  
Over18           0  
OverTime          0  
PercentSalaryHike 0  
PerformanceRating 0  
RelationshipSatisfaction 0  
StandardHours      0  
StockOptionLevel    0  
TotalWorkingYears   0  
TrainingTimesLastYear 0  
WorkLifeBalance    0  
YearsAtCompany     0  
YearsInCurrentRole 0  
YearsSinceLastPromotion 0  
YearsWithCurrManager 0  
dtype: int64
```

```
In [14]: df["Department"].value_counts()
```

```
Out[14]: Research & Development    961  
Sales                  446  
Human Resources        63  
Name: Department, dtype: int64
```

```
In [15]: df["BusinessTravel"].value_counts()
```

```
Out[15]: Travel_Rarely      1043  
Travel_Frequently      277  
Non-Travel            150  
Name: BusinessTravel, dtype: int64
```

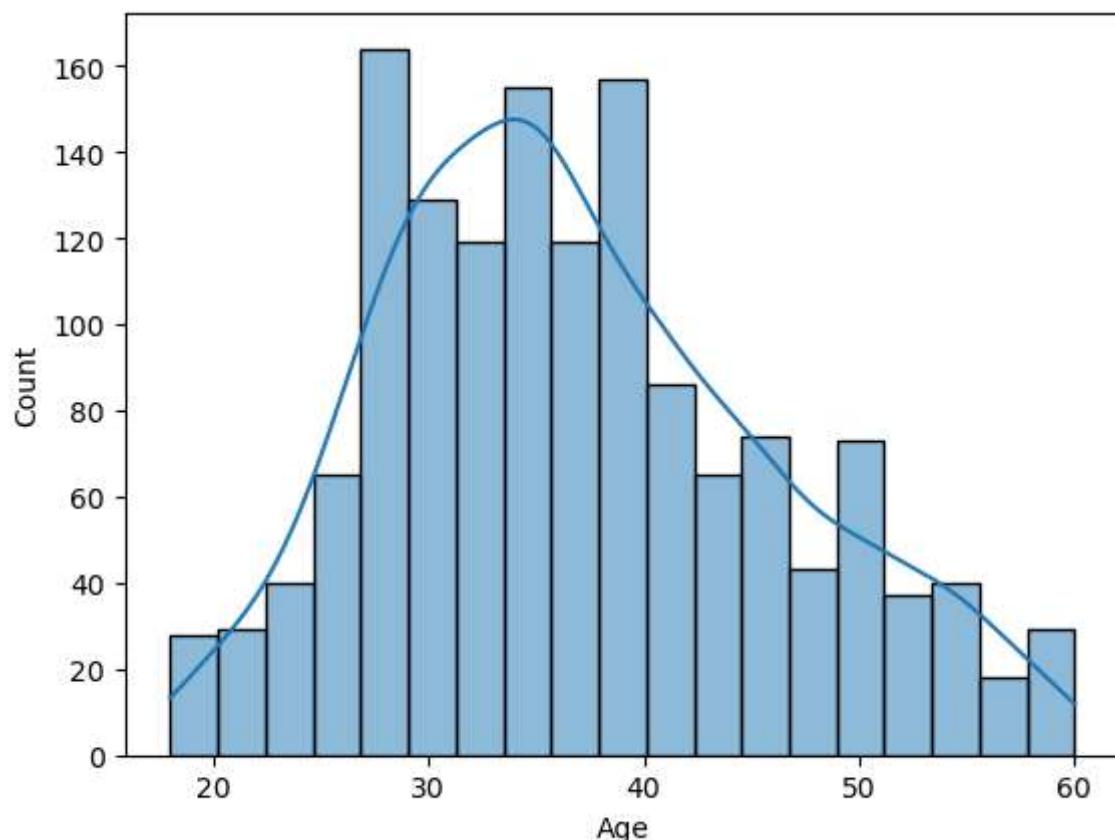
```
In [16]: df["EducationField"].value_counts()
```

```
Out[16]: Life Sciences      606  
Medical          464  
Marketing         159  
Technical Degree  132  
Other             82  
Human Resources   27  
Name: EducationField, dtype: int64
```

4.Data Visualization.

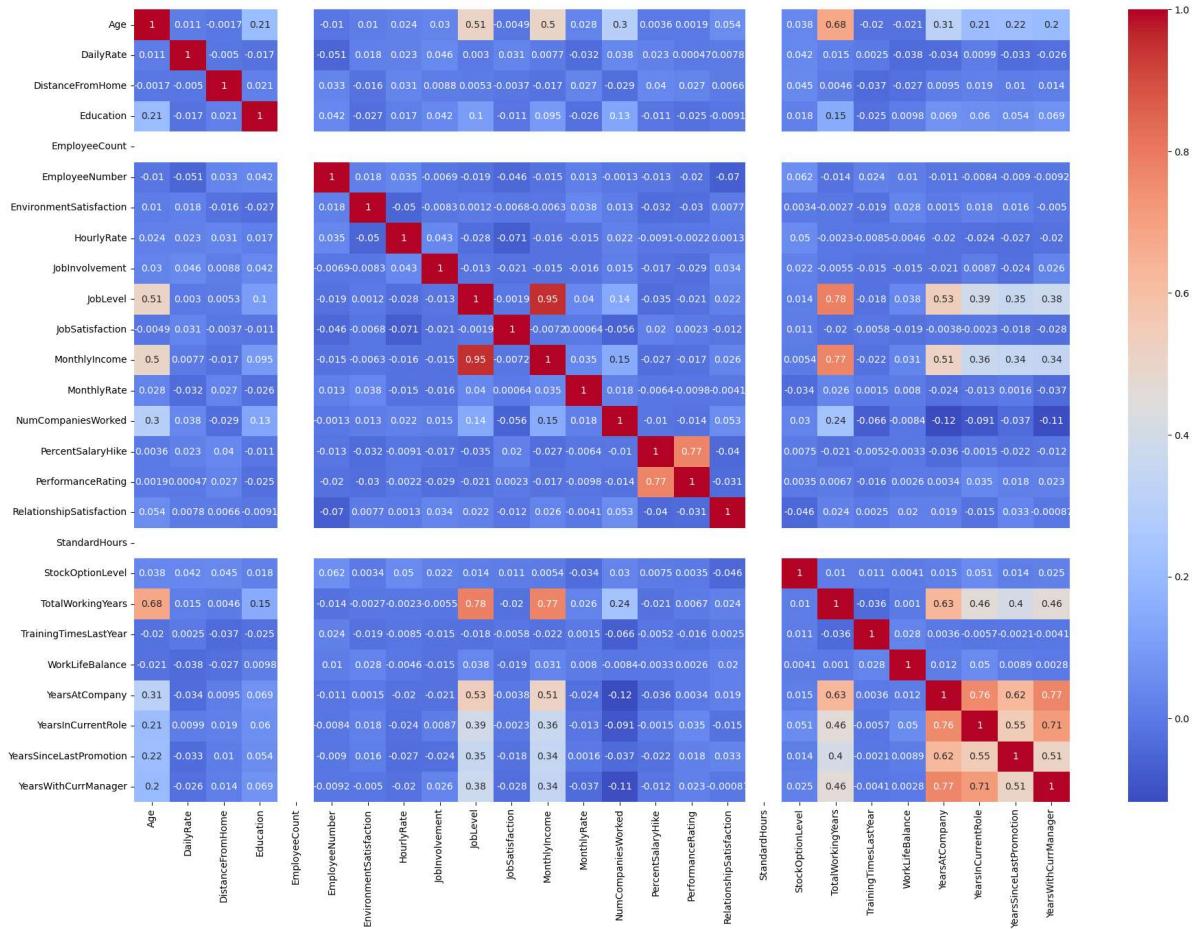
```
In [17]: sns.histplot(x="Age", data=df, kde=True)
```

```
Out[17]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



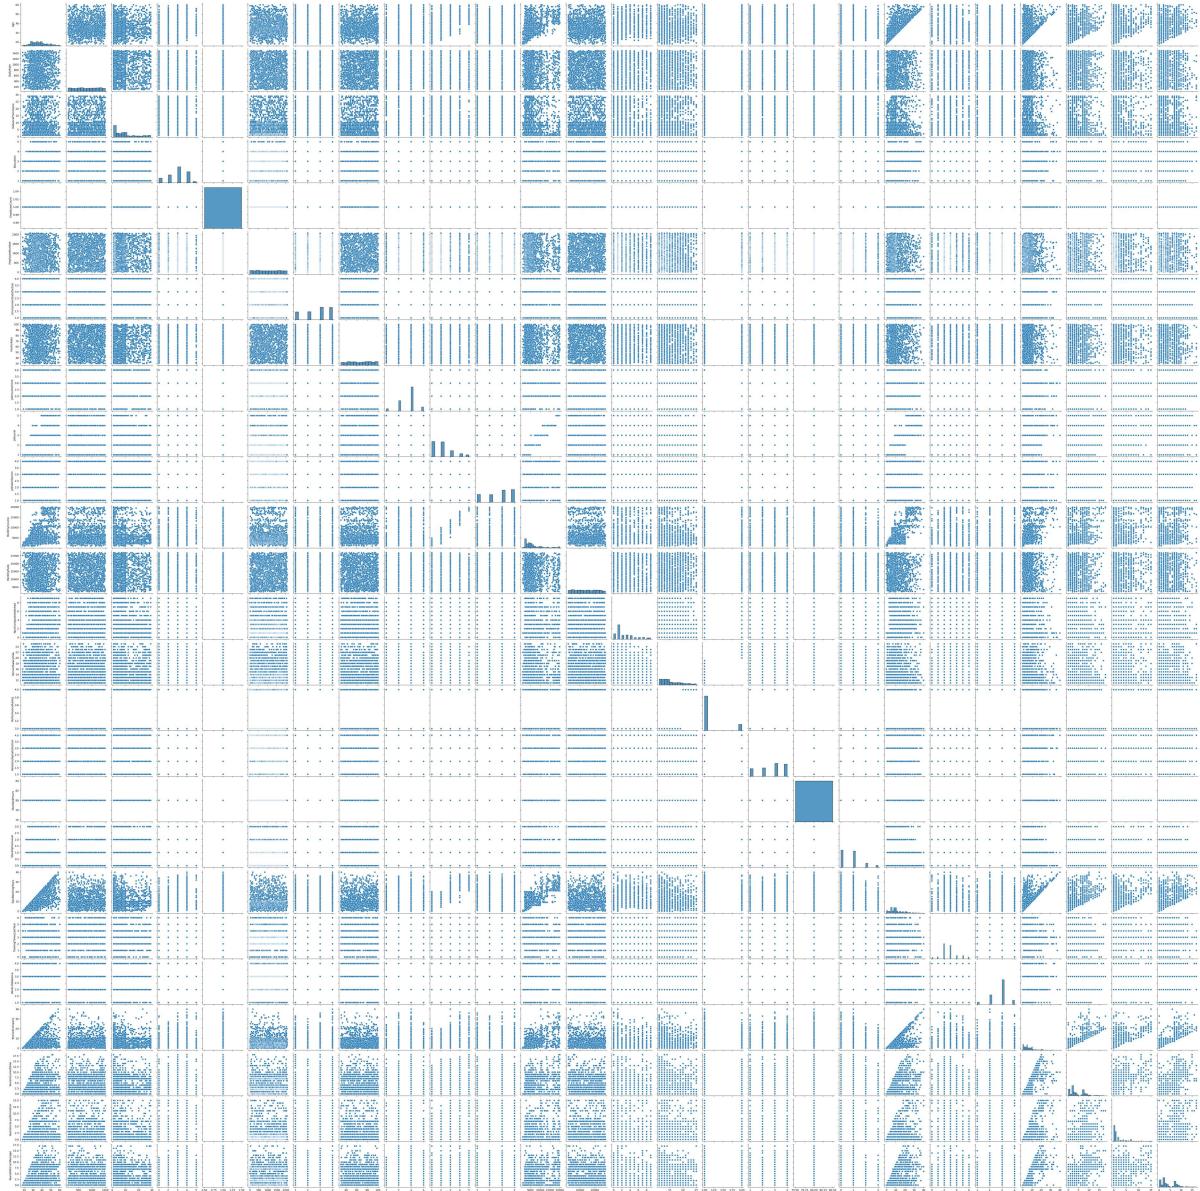
In [15]: `plt.subplots(figsize=(22,15))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm")`

Out[15]: <AxesSubplot:>



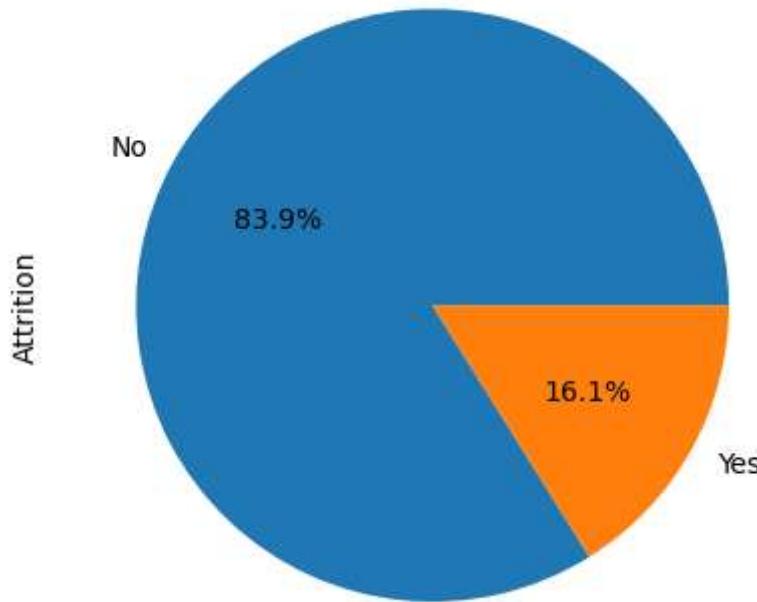
```
In [16]: sns.pairplot(df)
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x1e9dd3441f0>
```



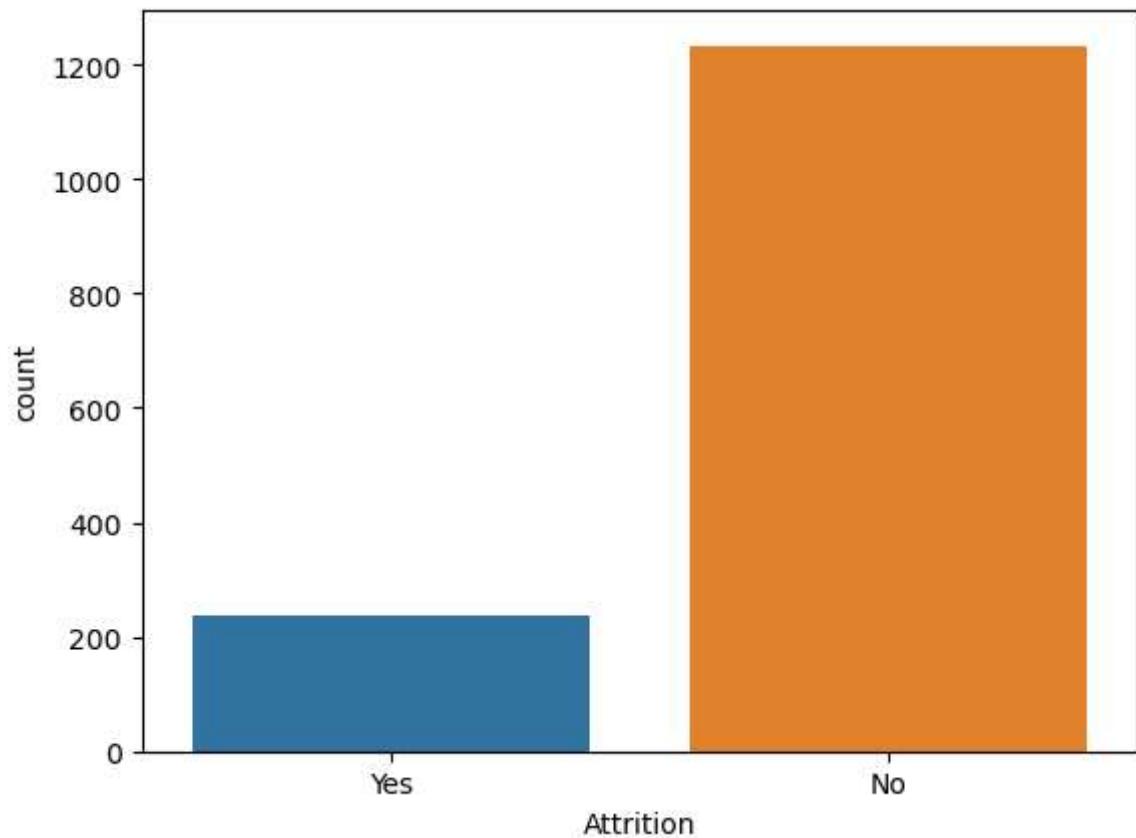
```
In [18]: df.Attrition.value_counts().plot(kind="pie", autopct="%1.1f%%")
```

```
Out[18]: <AxesSubplot:ylabel='Attrition'>
```



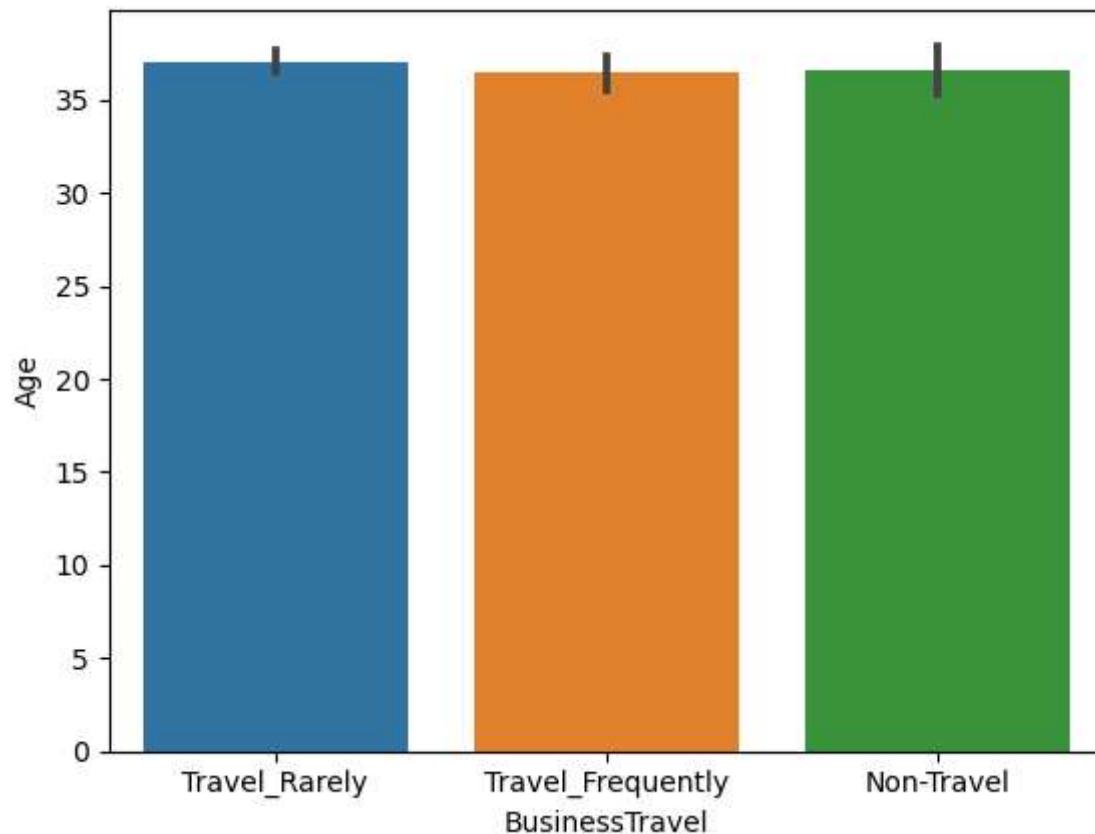
```
In [19]: sns.countplot(x="Attrition", data=df)
```

```
Out[19]: <AxesSubplot:xlabel='Attrition', ylabel='count'>
```



```
In [20]: sns.barplot(x=df["BusinessTravel"],y=df["Age"])
```

```
Out[20]: <AxesSubplot:xlabel='BusinessTravel', ylabel='Age'>
```



5. Outlier Detection

```
In [21]: df.head()
```

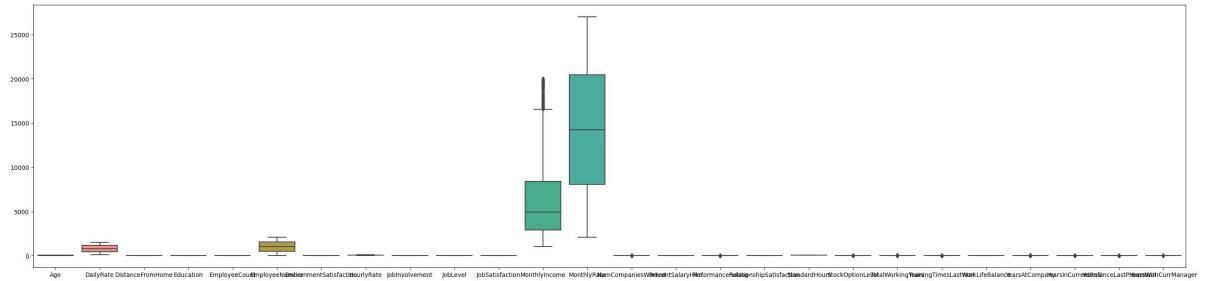
```
Out[21]:
```

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | Education |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|-----------|
| 0 | 41 | Yes | Travel_Rarely | 1102 | Sales | 1 | 2 | Life |
| 1 | 49 | No | Travel_Frequently | 279 | Research & Development | 8 | 1 | Life |
| 2 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | 2 | 2 | Life |
| 3 | 33 | No | Travel_Frequently | 1392 | Research & Development | 3 | 4 | Life |
| 4 | 27 | No | Travel_Rarely | 591 | Research & Development | 2 | 1 | Life |

5 rows × 35 columns

```
In [22]: plt.figure(figsize=(35,8))
sns.boxplot(data=df)
```

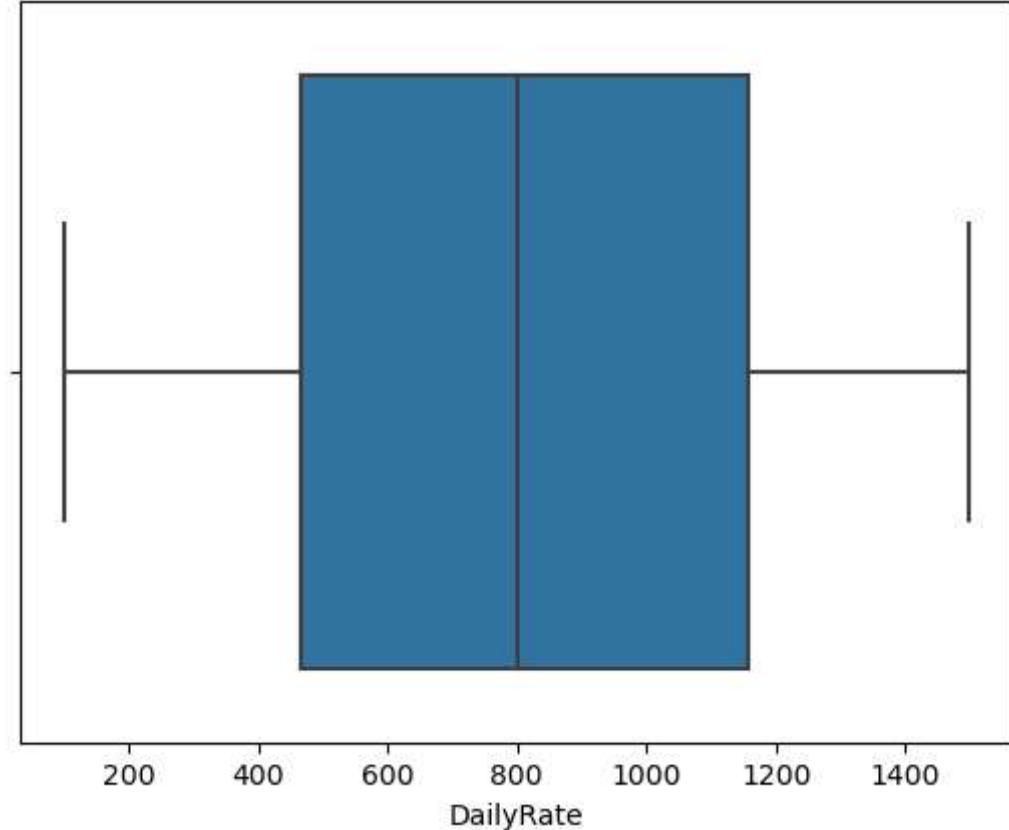
Out[22]: <AxesSubplot:>



```
In [23]: sns.boxplot(df[ "DailyRate" ])
```

C:\Users\admin\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

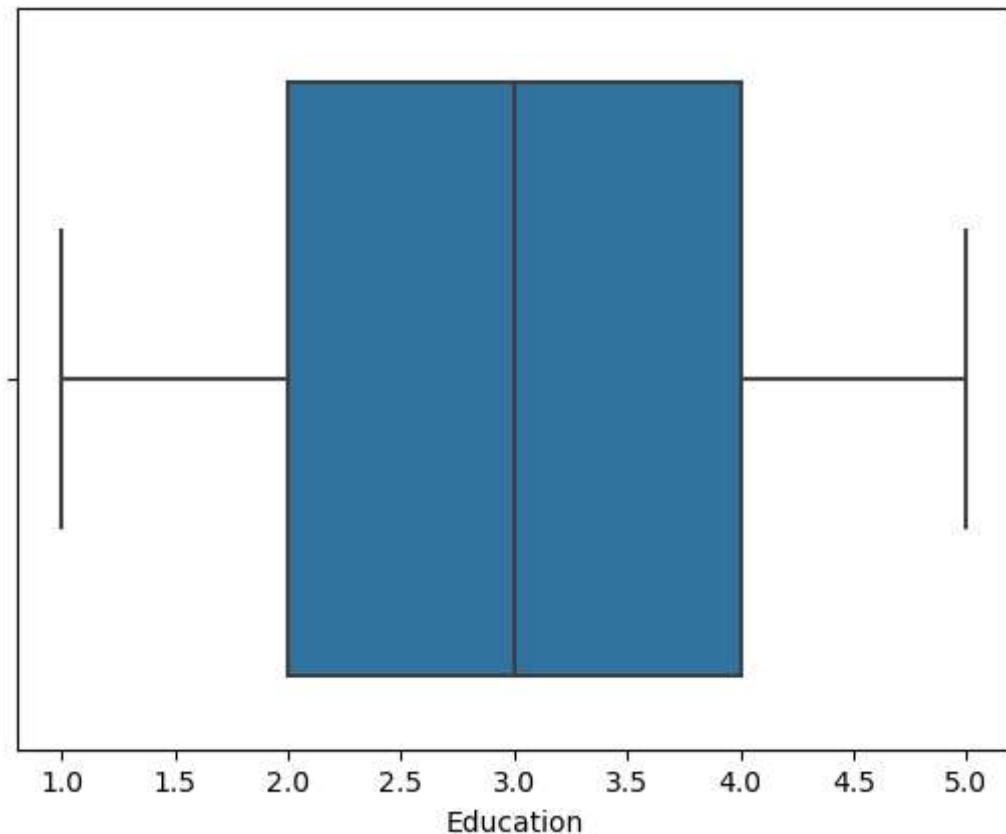
Out[23]: <AxesSubplot:xlabel='DailyRate'>



In [24]: `sns.boxplot(df["Education"])`

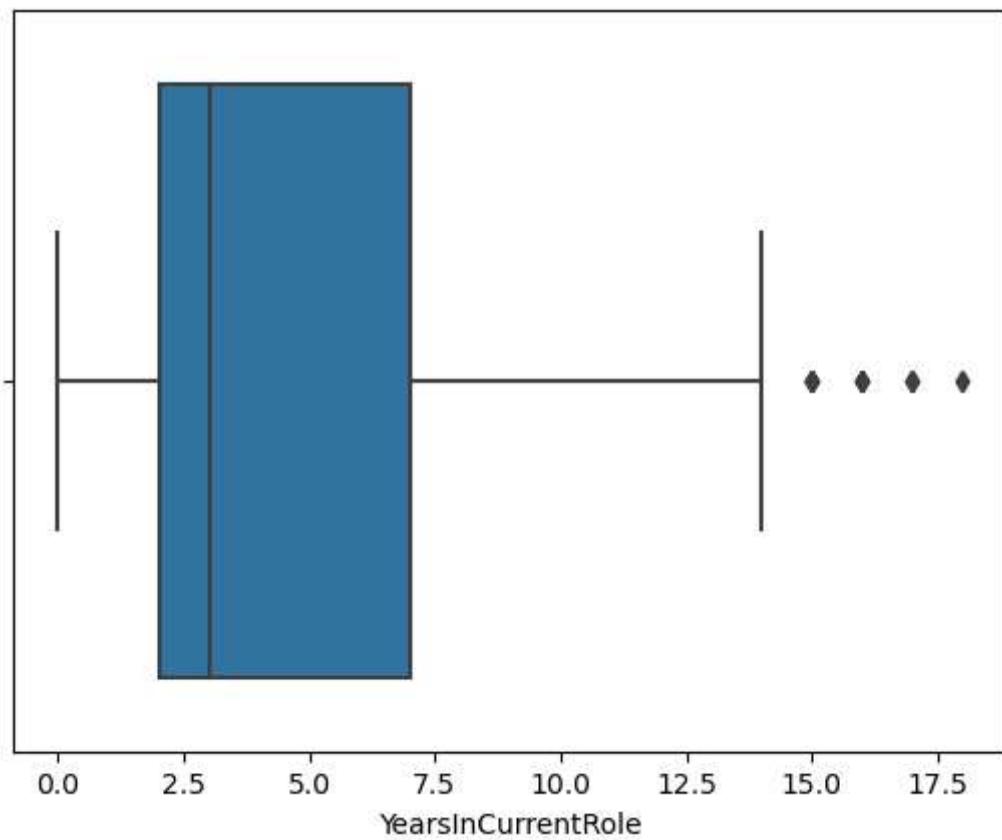
C:\Users\admin\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[24]: <AxesSubplot:xlabel='Education'>



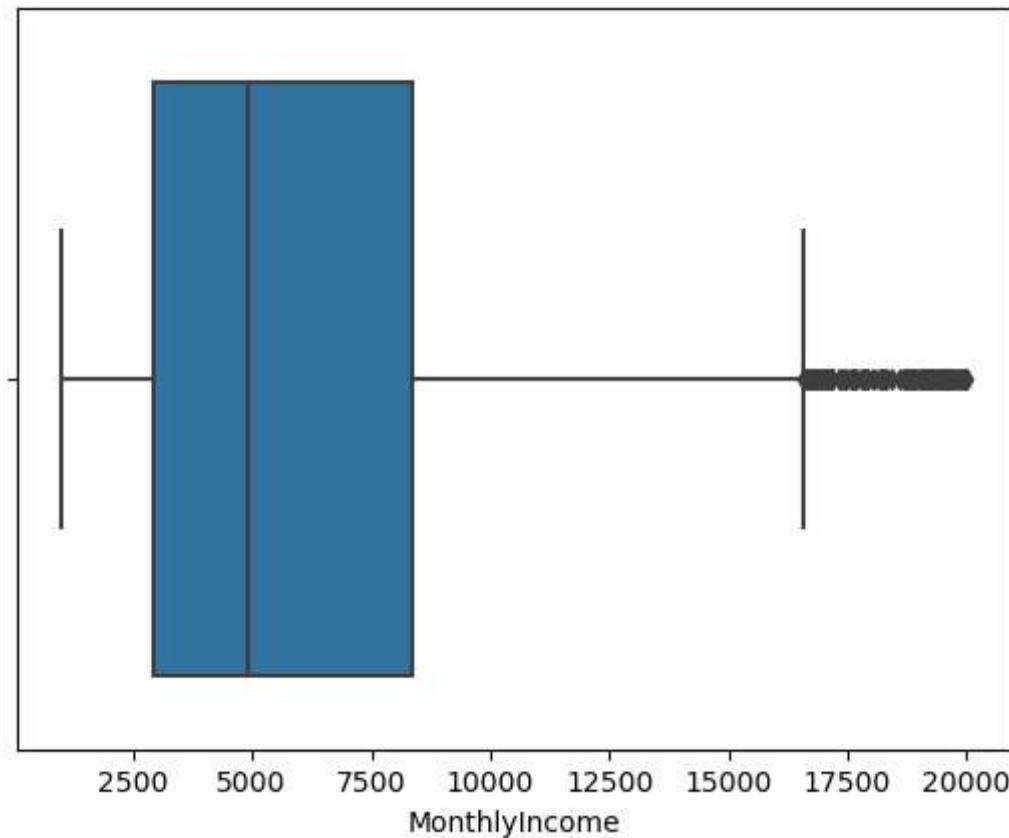
```
In [25]: sns.boxplot(data=df, x='YearsInCurrentRole')
```

```
Out[25]: <AxesSubplot:xlabel='YearsInCurrentRole'>
```



```
In [26]: sns.boxplot(data=df, x='MonthlyIncome')
```

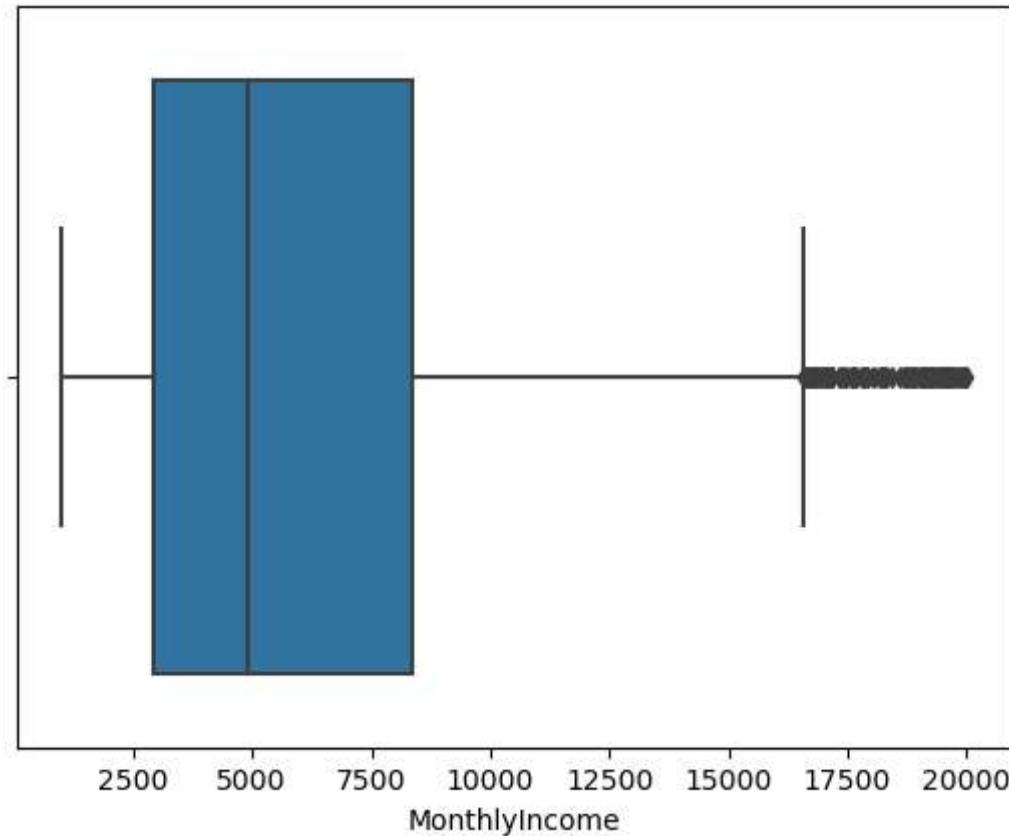
```
Out[26]: <AxesSubplot:xlabel='MonthlyIncome'>
```



```
In [27]: from scipy import stats  
z_scores = stats.zscore(df['MonthlyIncome'])  
df_cleaned = df[(np.abs(z_scores) <= 3)]
```

```
In [28]: sns.boxplot(data=df_cleaned, x='MonthlyIncome')
```

```
Out[28]: <AxesSubplot:xlabel='MonthlyIncome'>
```



So the outliers are in large quantity and they are inside the threshold, so let us not remove the outliers

6. Splitting Dependent and Independent variables

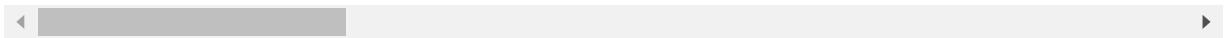
```
In [29]: x=df.drop(columns=["Attrition"],axis=1)
```

In [30]: `x.head()`

Out[30]:

| | Age | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | EducationField |
|---|-----|-------------------|-----------|------------------------|------------------|-----------|-----------------|
| 0 | 41 | Travel_Rarely | 1102 | Sales | | 1 | 2 Life Sciences |
| 1 | 49 | Travel_Frequently | 279 | Research & Development | | 8 | 1 Life Sciences |
| 2 | 37 | Travel_Rarely | 1373 | Research & Development | | 2 | 2 Other |
| 3 | 33 | Travel_Frequently | 1392 | Research & Development | | 3 | 4 Life Sciences |
| 4 | 27 | Travel_Rarely | 591 | Research & Development | | 2 | 1 Medical |

5 rows × 34 columns



In [31]: `type(x)`

Out[31]: `pandas.core.frame.DataFrame`

In [32]: `y=df["Attrition"]`

In [33]: `y.head()`

Out[33]:

| | |
|---|-----|
| 0 | Yes |
| 1 | No |
| 2 | Yes |
| 3 | No |
| 4 | No |

Name: Attrition, dtype: object

In [34]: `type(y)`

Out[34]: `pandas.core.series.Series`

7.Encoding

In [35]: `from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()`

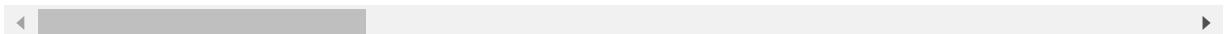
In [36]: `categorical_features = x.select_dtypes(include=['object']).columns.tolist()
x_encoded = pd.get_dummies(x, columns=categorical_features, drop_first=True)`

In [37]: `x_encoded.head()`

Out[37]:

| | Age | DailyRate | DistanceFromHome | Education | EmployeeCount | EmployeeNumber | EnvironmentSatisfaction | HourlyRate | JobInvolvement | JobLevel | JobRole | JobSatisfaction | MaritalStatus | OverTime | RelationshipSatisfaction | StandardHours | TotalWorkingYears | WorkLifeBalance | YearsAtCompany | YearsInCurrentRole | YearsOnSite |
|---|-----|-----------|------------------|-----------|---------------|----------------|-------------------------|------------|----------------|----------|---------|-----------------|---------------|----------|--------------------------|---------------|-------------------|-----------------|----------------|--------------------|-------------|
| 0 | 41 | 1102 | | 1 | 2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| 1 | 49 | 279 | | 8 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 2 | 1 | |
| 2 | 37 | 1373 | | 2 | 2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 4 | 1 | |
| 3 | 33 | 1392 | | 3 | 4 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 5 | 1 | |
| 4 | 27 | 591 | | 2 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 7 | 1 | |

5 rows × 47 columns



8.Feature Scaling

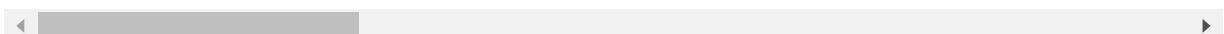
In [38]: `from sklearn.preprocessing import StandardScaler
s=StandardScaler()
x_scaled=pd.DataFrame(s.fit_transform(x_encoded),columns=x_encoded.columns)`

In [39]: `x_scaled.head()`

Out[39]:

| | Age | DailyRate | DistanceFromHome | Education | EmployeeCount | EmployeeNumber | EnvironmentSatisfaction | HourlyRate | JobInvolvement | JobLevel | JobRole | JobSatisfaction | MaritalStatus | OverTime | RelationshipSatisfaction | StandardHours | TotalWorkingYears | WorkLifeBalance | YearsAtCompany | YearsInCurrentRole | YearsOnSite |
|---|-----------|-----------|------------------|-----------|---------------|----------------|-------------------------|------------|----------------|----------|---------|-----------------|---------------|----------|--------------------------|---------------|-------------------|-----------------|----------------|--------------------|-------------|
| 0 | 0.446350 | 0.742527 | | -1.010909 | -0.891688 | | 0.0 | | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| 1 | 1.322365 | -1.297775 | | -0.147150 | -1.868426 | | 0.0 | | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| 2 | 0.008343 | 1.414363 | | -0.887515 | -0.891688 | | 0.0 | | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| 3 | -0.429664 | 1.461466 | | -0.764121 | 1.061787 | | 0.0 | | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |
| 4 | -1.086676 | -0.524295 | | -0.887515 | -1.868426 | | 0.0 | | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | |

5 rows × 47 columns



9.Splitting Data into Train and Test.

In [40]: `from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x_scaled,y,test_size = 0.2,random_state=42)`

In [41]: `x_train.shape,x_test.shape,y_train.shape,y_test.shape`

Out[41]: `((1176, 47), (294, 47), (1176,), (294,))`

• Model Building

- o Import the model building Libraries
- o Initializing the model
- o Training and testing the model
- o Evaluation of Model
- o Save the Model

Logistic Regression

1.Import the model building Libraries

```
In [42]: from sklearn.linear_model import LogisticRegression
```

2.Initializing the model

```
In [43]: lr=LogisticRegression()
```

3.Training and testing the model

```
In [44]: lr.fit(x_train,y_train)
```

```
Out[44]: LogisticRegression()
```

```
In [45]: y_pred=lr.predict(x_test)
```

In [96]: y_pred

In [97]: y_test

```
Out[97]:
```

| | |
|------|-----|
| 442 | No |
| 1091 | No |
| 981 | Yes |
| 785 | No |
| 1332 | Yes |
| | ... |
| 1439 | No |
| 481 | No |
| 124 | Yes |
| 198 | No |
| 1229 | No |

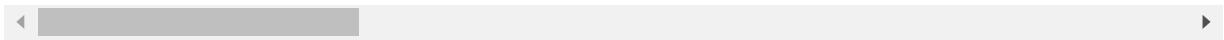
```
Name: Attrition, Length: 294, dtype: object
```

In [48]: `df.head()`

Out[48]:

| | Age | Attrition | BusinessTravel | DailyRate | Department | DistanceFromHome | Education | Education | |
|---|-----|-----------|-------------------|-----------|------------------------|------------------|-----------|-----------|------|
| 0 | 41 | Yes | Travel_Rarely | 1102 | Sales | | 1 | 2 | Life |
| 1 | 49 | No | Travel_Frequently | 279 | Research & Development | | 8 | 1 | Life |
| 2 | 37 | Yes | Travel_Rarely | 1373 | Research & Development | | 2 | 2 | Life |
| 3 | 33 | No | Travel_Frequently | 1392 | Research & Development | | 3 | 4 | Life |
| 4 | 27 | No | Travel_Rarely | 591 | Research & Development | | 2 | 1 | Life |

5 rows × 35 columns



4.Evaluation of Model

In [49]: `from sklearn.metrics import accuracy_score,confusion_matrix,classification_report`

In [50]: `accuracy_score(y_test,y_pred)`

Out[50]: `0.8775510204081632`

In [51]: `confusion_matrix(y_test,y_pred)`

Out[51]: `array([[238, 7], [29, 20]], dtype=int64)`

In [52]: `pd.crosstab(y_test,y_pred)`

Out[52]:

| col_0 | No | Yes |
|-------|----|-----|
|-------|----|-----|

Attrition

| | | |
|----|-----|---|
| No | 238 | 7 |
|----|-----|---|

| | | |
|-----|----|----|
| Yes | 29 | 20 |
|-----|----|----|

In [53]: `print(classification_report(y_test,y_pred))`

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| No | 0.89 | 0.97 | 0.93 | 245 |
| Yes | 0.74 | 0.41 | 0.53 | 49 |
| accuracy | | | 0.88 | 294 |
| macro avg | 0.82 | 0.69 | 0.73 | 294 |
| weighted avg | 0.87 | 0.88 | 0.86 | 294 |

```
In [54]: probability=lr.predict_proba(x_test)[:,1]
```

In [55]: probability

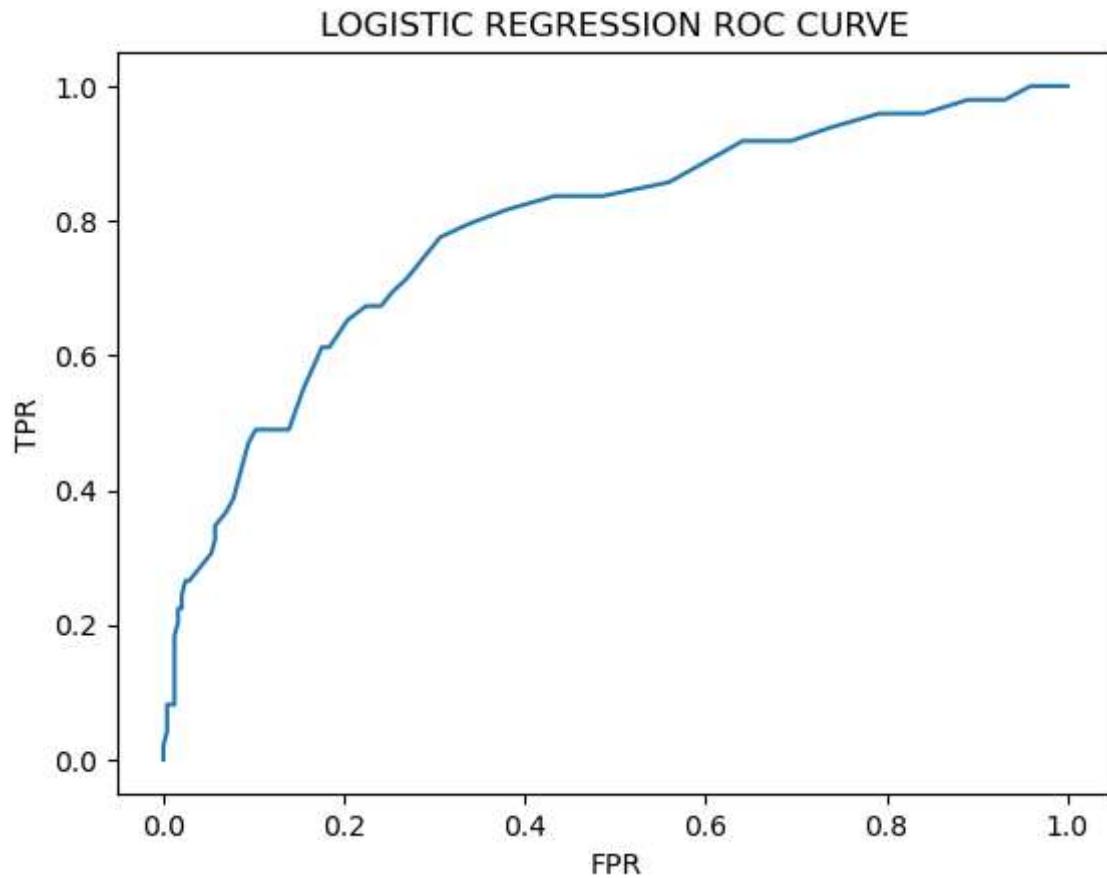
```
Out[55]: array([6.61367680e-02, 8.06198087e-02, 5.95640954e-01, 1.42633628e-01,
 7.84184536e-01, 4.42409704e-02, 4.90141163e-01, 3.64319554e-02,
 8.44925839e-04, 4.30164151e-01, 3.70048030e-02, 1.96283567e-01,
 1.51964558e-02, 5.39471444e-01, 5.88816107e-02, 1.18804547e-02,
 1.13489514e-01, 3.82574810e-02, 2.32011787e-02, 2.66149260e-01,
 1.31568762e-01, 9.99913247e-03, 1.30286201e-02, 3.88975456e-02,
 7.51582081e-01, 4.58920880e-01, 6.43869830e-02, 2.97280577e-02,
 6.90209855e-01, 3.77607239e-02, 5.59306642e-03, 3.05902665e-01,
 3.65028773e-02, 3.90288488e-02, 2.27899092e-02, 5.90358760e-03,
 2.32731579e-01, 5.29386349e-02, 1.38241033e-02, 1.13353905e-01,
 3.42777217e-02, 8.33820691e-03, 1.07151981e-03, 5.52578825e-03,
 8.80282261e-03, 5.59254552e-01, 4.21978392e-01, 5.99712843e-04,
 4.00447399e-01, 3.68527002e-01, 3.55627577e-02, 8.16716675e-01,
 1.58976261e-02, 3.55101528e-01, 4.94410470e-01, 2.54453615e-01,
 1.03880530e-02, 4.22173742e-01, 2.22235299e-02, 2.95631863e-01,
 1.26998325e-02, 1.53660325e-01, 1.65718568e-01, 2.37981819e-02,
 1.13922041e-01, 2.33156457e-02, 3.48112915e-01, 1.78088100e-01,
 1.21220858e-01, 1.82213934e-01, 3.38118048e-02, 1.58551253e-01,
 9.58027931e-02, 3.08386596e-02, 7.44773642e-02, 4.13402445e-02,
 1.92788840e-02, 2.80415152e-02, 4.93838215e-01, 1.65502835e-02,
 3.08277540e-03, 1.74647710e-02, 2.61191408e-01, 1.97665118e-02,
 1.95892874e-02, 7.95183218e-02, 6.72183840e-04, 1.32195797e-02,
 1.24777389e-02, 6.32250258e-02, 5.79625808e-02, 7.28350500e-02,
 3.24892323e-01, 1.73612641e-01, 8.28152031e-04, 7.99639374e-02,
 4.97002372e-01, 6.41780395e-01, 5.52704989e-02, 7.62263342e-02,
 1.32740233e-01, 5.87050953e-01, 5.91342152e-01, 5.08443347e-03,
 1.06560968e-01, 6.69002781e-03, 6.09068154e-02, 2.11328359e-01,
 3.68701854e-02, 3.96643551e-01, 2.94309518e-02, 5.18292700e-02,
 3.44618460e-03, 2.96577876e-01, 2.24618906e-02, 2.45843670e-02,
 9.84944236e-03, 2.82977886e-02, 2.36629032e-03, 4.89083971e-03,
 1.06593614e-01, 3.22277227e-02, 8.43908260e-02, 8.89755562e-01,
 1.01655842e-02, 2.07510840e-03, 3.14944475e-03, 9.91294127e-02,
 1.39271996e-01, 1.05763678e-02, 4.62491869e-03, 3.29517891e-01,
 7.00654301e-01, 1.96667235e-01, 9.97786091e-03, 3.85631608e-01,
 6.96550561e-01, 2.28516613e-01, 1.82557187e-01, 4.88667496e-01,
 2.73199468e-02, 3.53431498e-02, 3.17068584e-02, 3.13677515e-01,
 4.71792593e-01, 3.75527841e-02, 2.67239299e-01, 3.83012181e-03,
 7.37038139e-02, 1.24740256e-01, 4.07667992e-03, 1.71768746e-01,
 7.20048951e-02, 1.75320014e-01, 7.13178095e-03, 2.39429405e-02,
 1.98658569e-01, 3.20547275e-01, 3.96927966e-03, 9.49330798e-03,
 7.45555587e-01, 3.12330244e-03, 3.32408074e-02, 9.19392780e-01,
 1.48852250e-02, 2.57190744e-01, 1.73014525e-01, 1.51785386e-01,
 1.86931379e-02, 8.28851842e-04, 2.54342547e-01, 4.56483830e-02,
 3.99204597e-02, 1.25490112e-01, 1.70919482e-02, 8.25299901e-02,
 4.38992386e-02, 4.31506741e-02, 3.39152219e-02, 5.12964217e-02,
 2.12219051e-02, 1.57474151e-01, 2.44952859e-03, 8.22895146e-01,
 5.98438753e-02, 8.69783952e-02, 5.40701108e-01, 8.54091278e-03,
 5.57078600e-01, 2.31232882e-01, 2.11246138e-01, 2.75461703e-01,
 1.33154014e-01, 1.82199815e-02, 3.32865074e-02, 1.09429647e-01,
 1.86460970e-02, 7.81476677e-03, 2.65574308e-01, 1.69272765e-02,
 3.86686532e-01, 2.20988923e-01, 8.08612471e-01, 2.03966645e-02,
 1.28650474e-01, 1.77680007e-02, 3.57289398e-01, 5.30474583e-04,
 1.67105418e-01, 1.00764993e-02, 8.58290380e-02, 3.05739812e-01,
 6.81826517e-02, 3.81711531e-01, 1.10859745e-01, 4.71842280e-03,
 1.86894274e-02, 5.45048658e-02, 4.72421338e-02, 9.34489342e-02,
 8.33653672e-02, 4.44029727e-01, 5.18663398e-01, 1.79271982e-01,
 2.87090882e-01, 2.67710024e-03, 1.09509162e-01, 2.65241921e-01,
```

```
7.65032998e-01, 4.64180665e-02, 1.97677279e-01, 2.06495230e-01,  
3.00475823e-02, 4.28680929e-02, 9.01257179e-02, 1.16204269e-01,  
2.66591825e-01, 9.07332777e-04, 7.47052094e-02, 1.49451887e-03,  
1.41520152e-01, 2.92363790e-01, 5.07029822e-03, 1.26501783e-01,  
3.79539275e-02, 1.77150306e-02, 1.54456455e-01, 3.43384106e-01,  
4.27481928e-02, 3.44423944e-02, 2.93165700e-01, 7.16660273e-02,  
5.38151329e-01, 1.63749013e-02, 1.23951727e-01, 6.52278702e-02,  
2.45783680e-03, 6.25002778e-01, 6.25608255e-01, 3.74616723e-01,  
1.41256456e-01, 2.84604841e-02, 2.91843045e-01, 5.95069780e-02,  
3.21045151e-02, 6.99486994e-02, 1.15155181e-03, 4.47317458e-01,  
5.39084197e-01, 2.55697184e-02, 2.85616158e-02, 8.43094503e-03,  
6.55595901e-02, 2.48427612e-02, 9.60265542e-03, 5.93422188e-03,  
8.58521685e-02, 3.64680251e-01, 1.20449475e-01, 1.84226608e-01,  
6.91761349e-01, 3.24455148e-03, 4.85788438e-02, 7.34675868e-02,  
5.35355633e-02, 7.58287245e-02, 2.30602991e-04, 1.53563842e-01,  
1.43056931e-03, 9.15060156e-03, 1.08956580e-01, 6.50047900e-01,  
2.36548949e-02, 7.52606168e-02])
```

```
In [98]: from sklearn.preprocessing import LabelEncoder  
label_encoder = LabelEncoder()  
y_test = label_encoder.fit_transform(y_test)
```

```
In [99]: fpr,tpr,thresholds = roc_curve(y_test,probability)
```

```
In [100]: plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('LOGISTIC REGRESSION ROC CURVE')
plt.show()
```



Decision Tree

```
In [57]: from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```

```
In [58]: dtc.fit(x_train,y_train)
```

```
Out[58]: DecisionTreeClassifier()
```

```
In [59]: y_pred1=dtc.predict(x_test)
```

In [101]: y_pred1

```
Out[101]: array(['No', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No',
       'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
       'Yes', 'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No',
       'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
       'Yes', 'No', 'No', 'No', 'Yes', 'No', 'Yes', 'No', 'No', 'Yes',
       'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No',
       'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
       'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No'],
      dtype=object)
```

In [64]: y_test

```
Out[64]: 442      No
1091     No
981      Yes
785      No
1332     Yes
...
1439     No
481      No
124      Yes
198      No
1229     No
Name: Attrition, Length: 294, dtype: object
```

In [65]: #Accuracy score

```
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

In [66]: accuracy_score(y_test,y_pred1)

```
Out[66]: 0.7687074829931972
```

```
In [67]: confusion_matrix(y_test,y_pred1)
```

```
Out[67]: array([[210,  35],  
                 [ 33,  16]], dtype=int64)
```

```
In [68]: pd.crosstab(y_test,y_pred1)
```

Out[68]:

col 0 No Yes

Attrition

| | | |
|-----|-----|----|
| No | 210 | 35 |
| Yes | 33 | 16 |

Yes 33 16

```
In [69]: print(classification_report(y_test,y_pred1))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| No | 0.86 | 0.86 | 0.86 | 245 |
| Yes | 0.31 | 0.33 | 0.32 | 49 |
| accuracy | | | 0.77 | 294 |
| macro avg | 0.59 | 0.59 | 0.59 | 294 |
| weighted avg | 0.77 | 0.77 | 0.77 | 294 |

```
In [70]: probability=dtc.predict_proba(x_test)[:,1]
```

```
In [71]: probability
```

Hyper parameter Tuning

```
In [67]: from sklearn import tree
plt.figure(figsize=(25,15))
tree.plot_tree(dtc,filled=True)
```

```
Out[67]: [Text(0.32960321576763485, 0.96875, 'X[19] <= -1.257\ngini = 0.269\nsamples = 1176\nvalue = [988, 188]'),
Text(0.08298755186721991, 0.90625, 'X[45] <= 0.387\ngini = 0.5\nsamples = 78\nvalue = [39, 39]'),
Text(0.04979253112033195, 0.84375, 'X[2] <= 0.902\ngini = 0.426\nsamples = 39\nvalue = [27, 12]'),
Text(0.03319502074688797, 0.78125, 'X[26] <= 0.797\ngini = 0.312\nsamples = 31\nvalue = [25, 6]'),
Text(0.01991701244813278, 0.71875, 'X[10] <= -1.114\ngini = 0.198\nsamples = 27\nvalue = [24, 3]'),
Text(0.013278008298755186, 0.65625, 'X[46] <= 0.482\ngini = 0.5\nsamples = 6\nvalue = [3, 3]'),
Text(0.006639004149377593, 0.59375, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.01991701244813278, 0.59375, 'gini = 0.0\nsamples = 3\nvalue = [0, 3]'),
Text(0.026556016597510373, 0.65625, 'gini = 0.0\nsamples = 21\nvalue = [21, 0]'),
Text(0.046473029045643155, 0.71875, 'X[6] <= 0.712\ngini = 0.375\nsamples = 11\nvalue = [11, 11]')]
```

```
In [72]: from sklearn.model_selection import GridSearchCV
parameter={
    'criterion':['gini','entropy'],
    'splitter':['best','random'],
    'max_depth':[1,2,3,4,5],
    'max_features':['auto', 'sqrt', 'log2']
}
```

```
In [73]: grid_search=GridSearchCV(estimator=dtc,param_grid=parameter,cv=5,scoring="accuracy")
```

```
In [74]: grid_search.fit(x_train,y_train)
```

```
Out[74]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
param_grid={'criterion': ['gini', 'entropy'],
'max_depth': [1, 2, 3, 4, 5],
'max_features': ['auto', 'sqrt', 'log2'],
'splitter': ['best', 'random']},
scoring='accuracy')
```

```
In [75]: grid_search.best_params_
```

```
Out[75]: {'criterion': 'entropy',
'max_depth': 4,
'max_features': 'auto',
'splitter': 'best'}
```

```
In [76]: dtc_cv=DecisionTreeClassifier(criterion= 'entropy',
max_depth=3,
max_features='sqrt',
splitter='best')
dtc_cv.fit(x_train,y_train)
```

```
Out[76]: DecisionTreeClassifier(criterion='entropy', max_depth=3, max_features='sqrt')
```

```
In [77]: y_pred=dtc_cv.predict(x_test)
```

```
In [78]: print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| No | 0.84 | 0.98 | 0.90 | 245 |
| Yes | 0.33 | 0.04 | 0.07 | 49 |
| accuracy | | | 0.83 | 294 |
| macro avg | 0.59 | 0.51 | 0.49 | 294 |
| weighted avg | 0.75 | 0.83 | 0.77 | 294 |

Random Forest

```
In [79]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
```

```
In [80]: rfc.fit(x_train,y_train)
```

```
Out[80]: RandomForestClassifier()
```

```
In [81]: y_pred2=dtc.predict(x_test)
```

In [82]: y_pred2

```
Out[82]: array(['No', 'No', 'Yes', 'No', 'Yes', 'Yes', 'No', 'No', 'No',
   'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No',
   'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No',
   'Yes', 'Yes', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No',
   'No', 'Yes', 'No', 'No', 'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No',
   'Yes', 'No', 'No', 'Yes', 'Yes', 'No', 'Yes', 'No', 'No', 'Yes', 'No',
   'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No', 'No',
   'No', 'Yes', 'No', 'No', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'No', 'No', 'No', 'No', 'No',
   'No', 'No', 'No', 'No', 'No', 'No', 'Yes', 'Yes', 'No', 'No', 'No', 'No'],
  dtype=object)
```

In [83]: y_test

```
Out[83]: 442      No
1091     No
981      Yes
785      No
1332     Yes
...
1439     No
481      No
124      Yes
198      No
1229     No
Name: Attrition, Length: 294, dtype: object
```

In [84]: forest_params = [{`'max_depth'`: list(range(10, 15)), `'max_features'`: list(range

In [85]: rfc_cv=GridSearchCV(rfc,param_grid=forest_params,cv=10,scoring='accuracy')

```
In [86]: rfc_cv.fit(x_train,y_train)
```

```
C:\Users\admin\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:372: FitFailedWarning:  
50 fits failed out of a total of 700.  
The score on these train-test partitions for these parameters will be set to  
nan.  
If these failures are not expected, you can try to debug them by setting erro  
r_score='raise'.
```

Below are more details about the failures:

```
-----  
---  
50 fits failed with the following error:  
Traceback (most recent call last):  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\model_selection\_v  
alidation.py", line 680, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\ensemble\_forest.p  
y", line 450, in fit  
    trees = Parallel(  
      File "C:\Users\admin\anaconda3\lib\site-packages\joblib\parallel.py", line  
1863, in __call__  
    return output if self.return_generator else list(output)  
  File "C:\Users\admin\anaconda3\lib\site-packages\joblib\parallel.py", line  
1792, in _get_sequential_output  
    res = func(*args, **kwargs)  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\utils\fixes.py", l  
ine 216, in __call__  
    return self.function(*args, **kwargs)  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\ensemble\_forest.p  
y", line 185, in _parallel_build_trees  
    tree.fit(X, y, sample_weight=curr_sample_weight, check_input=False)  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\tree\classes.py",  
line 937, in fit  
    super().fit()  
  File "C:\Users\admin\anaconda3\lib\site-packages\sklearn\tree\classes.py",  
line 308, in fit  
    raise ValueError("max_features must be in (0, n_features]")  
ValueError: max_features must be in (0, n_features]  
  
warnings.warn(some_fits_failed_message, FitFailedWarning)  
C:\Users\admin\anaconda3\lib\site-packages\sklearn\model_selection\_search.p  
y:969: UserWarning: One or more of the test scores are non-finite: [       na  
n 0.84948573 0.84950022 0.85034043 0.85459945 0.85801101  
 0.85628712 0.8605389 0.85885122 0.85798204 0.86052441 0.8596842  
 0.86138635 0.85796031       nan 0.84693611 0.85290453 0.85714907  
 0.85882949 0.85884398 0.85799652 0.85798204 0.85543242 0.8596842  
 0.8596842 0.85883674 0.8613936 0.85542518       nan 0.84948573  
 0.85204983 0.85630161 0.85798928 0.85628712 0.85543966 0.86222657  
 0.85883674 0.86309576 0.86053165 0.86223381 0.86308127 0.85797479  
       nan 0.84779806 0.85204983 0.85800377 0.85883674 0.86055338  
 0.85800377 0.86052441 0.8605389 0.85712009 0.85797479 0.85798928  
 0.86053165 0.86137187       nan 0.84863827 0.85545415 0.85629436  
 0.85884398 0.85798928 0.85373751 0.85799652 0.85799652 0.85712734  
 0.86140808 0.85712734 0.86222657 0.85967695]  
  warnings.warn(
```

```
Out[86]: GridSearchCV(cv=10, estimator=RandomForestClassifier(),
                      param_grid=[{'max_depth': [10, 11, 12, 13, 14],
                                   'max_features': [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10,
                                   11,
                                   12, 13]}],
                      scoring='accuracy')
```

```
In [87]: y_pred=rfc_cv.predict(x_test)
```

```
In [88]: from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
```

```
In [89]: accuracy_score(y_test,y_pred2)
```

```
Out[89]: 0.7687074829931972
```

```
In [90]: confusion_matrix(y_test,y_pred2)
```

```
Out[90]: array([[210, 35],
                 [ 33, 16]], dtype=int64)
```

```
In [91]: pd.crosstab(y_test,y_pred2)
```

```
Out[91]:
      col_0  No  Yes
Attrition
_____
No    210   35
Yes    33   16
```

```
In [92]: print(classification_report(y_test,y_pred2))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| No | 0.86 | 0.86 | 0.86 | 245 |
| Yes | 0.31 | 0.33 | 0.32 | 49 |
| accuracy | | | 0.77 | 294 |
| macro avg | 0.59 | 0.59 | 0.59 | 294 |
| weighted avg | 0.77 | 0.77 | 0.77 | 294 |

```
In [93]: rfc_cv.best_params_
```

```
Out[93]: {'max_depth': 12, 'max_features': 9}
```

```
In [94]: probability=rfc.predict_proba(x_test)[:,1]
```

In [95]: probability

Out[95]: array([0.13, 0.05, 0.18, 0.15, 0.71, 0.36, 0.3 , 0.11, 0.11, 0.13, 0.06, 0.11, 0.07, 0.47, 0.1 , 0.02, 0.16, 0.08, 0.11, 0.16, 0.56, 0.07, 0.05, 0.11, 0.39, 0.21, 0.06, 0.12, 0.57, 0.07, 0.06, 0.13, 0.13, 0.16, 0.07, 0.09, 0.16, 0.14, 0.03, 0.25, 0.19, 0.08, 0.11, 0.2 , 0.05, 0.36, 0.31, 0.08, 0.65, 0.33, 0.26, 0.48, 0.14, 0.21, 0.41, 0.14, 0.14, 0.1 , 0.05, 0.26, 0.03, 0.16, 0.11, 0.15, 0.37, 0.12, 0.21, 0.14, 0.09, 0.16, 0.1 , 0.3 , 0.12, 0.09, 0.1 , 0.11, 0.1 , 0.04, 0.39, 0.05, 0.05, 0.11, 0.23, 0.1 , 0.18, 0.08, 0.09, 0.25, 0.02, 0.04, 0.65, 0.1 , 0.13, 0.26, 0.06, 0.07, 0.1 , 0.32, 0.14, 0.13, 0.27, 0.24, 0.36, 0.05, 0.07, 0.03, 0.09, 0.47, 0.23, 0.27, 0.2 , 0.12, 0.04, 0.07, 0.15, 0.16, 0.1 , 0.19, 0.04, 0.01, 0.11, 0.03, 0.04, 0.58, 0.12, 0.16, 0.02, 0.06, 0.11, 0.12, 0.01, 0.29, 0.5 , 0.21, 0.18, 0.14, 0.35, 0.2 , 0.14, 0.16, 0.1 , 0.24, 0.09, 0.15, 0.24, 0.17, 0.06, 0.05, 0.15, 0.15, 0.1 , 0.23, 0.13, 0.23, 0.15, 0.12, 0.11, 0.21, 0.12, 0.09, 0.23, 0.07, 0.28, 0.51, 0.1 , 0.14, 0.23, 0.06, 0.11, 0.04, 0.12, 0.05, 0.1 , 0.25, 0.07, 0.43, 0.12, 0.25, 0.21, 0.09, 0.04, 0.02, 0.01, 0.37, 0.07, 0.03, 0.37, 0.07, 0.2 , 0.13, 0.24, 0.5 , 0.2 , 0.04, 0.1 , 0.08, 0.06, 0.04, 0.46, 0.12, 0.29, 0.19, 0.16, 0.04, 0.11, 0.17, 0.35, 0.07, 0.08, 0.01, 0.1 , 0.16, 0.09, 0.23, 0.1 , 0.08, 0.06, 0.25, 0.11, 0.37, 0.17, 0.36, 0.38, 0.17, 0.11, 0.08, 0.13, 0.27, 0.52, 0.12, 0.03, 0.31, 0.05, 0.08, 0.05, 0.26, 0.28, 0.03, 0.08, 0.1 , 0.21, 0.13, 0.1 , 0.13, 0.15, 0.09, 0.05, 0.32, 0.06, 0.15, 0.2 , 0.14, 0.24, 0.07, 0.18, 0.14, 0.08, 0.61, 0.23, 0.31, 0.16, 0.08, 0.13, 0.11, 0.1 , 0.22, 0.06, 0.28, 0.24, 0.03, 0.06, 0.02, 0.09, 0.1 , 0.07, 0.05, 0.13, 0.32, 0.09, 0.19, 0.32, 0.1 , 0.11, 0.24, 0.17, 0.22, 0.01, 0.1 , 0.1 , 0.12, 0.13, 0.28, 0.09, 0.11])

```
In [105]: # Import the necessary Libraries
from sklearn.metrics import roc_curve, roc_auc_score
import matplotlib.pyplot as plt

logreg_roc_auc = roc_auc_score(y_test, lr.predict_proba(x_test)[:, 1])

# Calculate the ROC AUC score for Decision Tree Classifier
dt_roc_auc = roc_auc_score(y_test, grid_search.predict_proba(x_test)[:, 1])

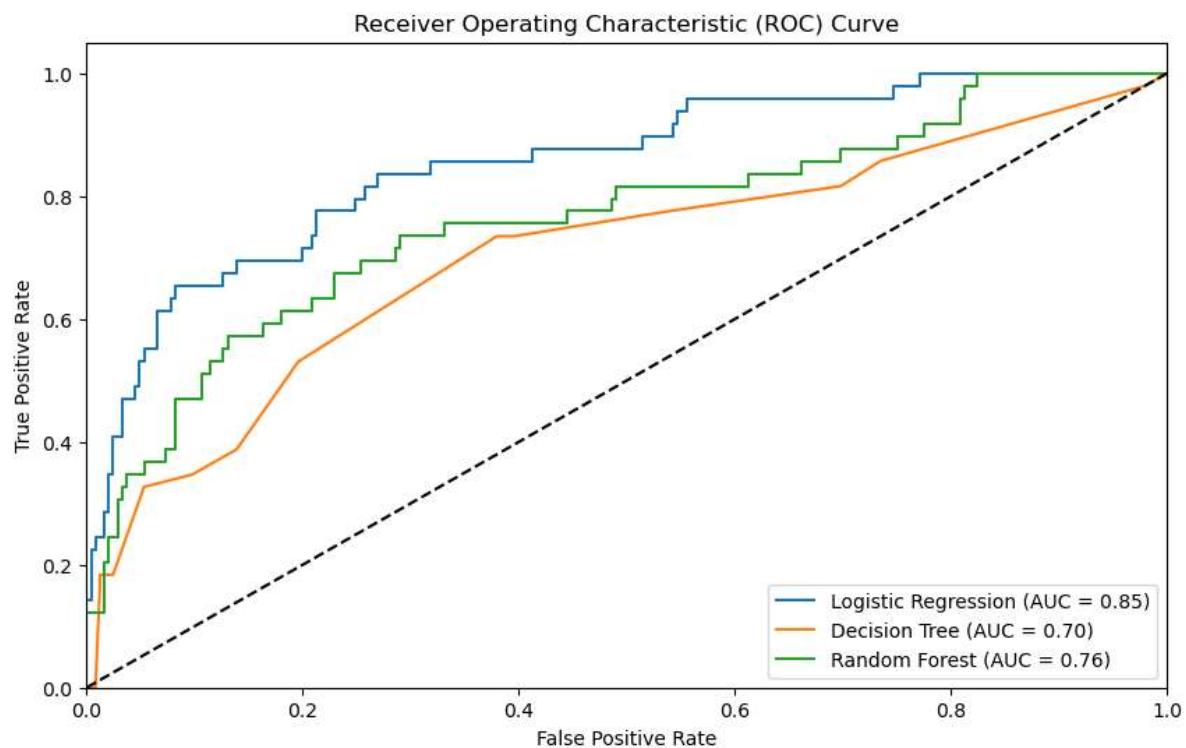
# Compute ROC curve for Logistic Regression
logreg_fpr, logreg_tpr, _ = roc_curve(y_test, lr.predict_proba(x_test)[:, 1])

# Compute ROC curve for Decision Tree Classifier
dt_fpr, dt_tpr, _ = roc_curve(y_test, grid_search.predict_proba(x_test)[:, 1])

# Calculate the ROC AUC score for Random Forest
rf_roc_auc = roc_auc_score(y_test, rfc_cv.predict_proba(x_test)[:, 1])

# Compute ROC curve for Random Forest
rf_fpr, rf_tpr, _ = roc_curve(y_test, rfc_cv.predict_proba(x_test)[:, 1])

# Plot ROC AUC curves
plt.figure(figsize=(10, 6))
plt.plot(logreg_fpr, logreg_tpr, label='Logistic Regression (AUC = %0.2f)' % logreg_roc_auc)
plt.plot(dt_fpr, dt_tpr, label='Decision Tree (AUC = %0.2f)' % dt_roc_auc)
plt.plot(rf_fpr, rf_tpr, label='Random Forest (AUC = %0.2f)' % rf_roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```



In []: